


(a) Representation-consistent output**(b) Masked post-adaptation (AMS)**

Training: new mask per image and epoch




Fig. 3. Representation-consistent output reconstruction and masked post-adaptation (AMS). (a) A bounded log prediction is exactly inverted before residual compensation and addition in the intensity domain. (b) During adaptation, the whole model receives the masked observation Y_m , including the compensator; the original observation supplies the noisy target. Only the Transformer encoder is frozen. Training masks are resampled, validation masks are fixed for checkpoint selection, and inference uses the complete unmasked observation.

Single-column figure

Size: 88.568 x 182.033 mm.

Keep width=columnwidth and natural height. The original 1.25-inch placeholder cannot retain readable labels at this level of detail.

Panel (a): numerical domain

The bounded log prediction is analytically inverted. Both the residual base and the compensator's inputs are intensities. Clipping follows the residual addition.

Panel (b): training and inference

The entire model receives Y_m during adaptation, including its observation bypass. Original Y supplies only the noisy loss target. M selects loss pixels. The shaded parameter groups show optimization status, not serial forward stages.

Validation reuses a fixed mask for each image and compares the same masked L1 plus TV objective. Inference uses the selected model without a mask.

Paper/code difference recorded

The paper writes $\text{sum}(M) + \epsilon$ in the masked-L1 denominator; the actual code uses $\max(\text{sum}(M), 1)$. The diagram names the implemented masked objective; SOURCE_AUDIT.md records the exact evidence. Neither code nor paper was altered.

Verified outputs

88 passed checks, including two full-model forwards.

PDF: vectors and embedded fonts.

SVG: 16 grouped modules; main text remains editable. Ten special math glyphs are portable vector outlines.

Scope

This proof verifies final figure size, not full-manuscript pagination. No experimental data, checkpoint, training, or model edit was used for the figure.